

A Polyphonic Conception of AI Understanding

Abstract: Whether LLMs understand what they talk about is a question that is both unavoidable and ill-framed. It is unavoidable because users must distinguish outputs that merit epistemic trust from those that do not, and mathematical or statistical descriptions cannot draw that distinction without reintroducing the concept of understanding in all but name. Yet the question is ill-framed as it stands, because the inherited concept operates within a *monophonic* paradigm, inviting us to look for one locus of understanding supporting application, explanation, prediction, and counterfactual reasoning alike. Mechanistic interpretability instead reveals LLMs to be deeply *polyphonic* systems whose outputs emerge from parallel mechanisms of uneven reliability. Drawing on mechanistic evidence of medical, mathematical, logical, and spatial understanding in LLMs, we engineer a conception of understanding fit for polyphonic AI systems. On this conception, a model understands a domain insofar as it contains *sound circuitry* – features, connections, and procedures that track relevant structure – and reliably recruits that circuitry to control its outputs. This turns seemingly idle metaphysical disputes into tractable empirical questions about the presence, recruitment, control, and scope of internal organisation, and does necessary work in guiding the allocation of trust in AI outputs.

Keywords: AI understanding; machine understanding; large language models; mechanistic interpretability; explainability; conceptual engineering; trustworthy AI; polyphony

1. The Jury Room Analogy

Imagine a jury room in which deliberation is conducted in relays, with a fresh panel of jurors coming in every morning. There is no presiding juror to orchestrate the proceedings. All that persist are the notes which each departing panel leaves on a large communal table. Incoming jurors read this material selectively and with varying degrees of competence. Some happen to have occupational expertise they can bring to bear; others reason more superficially, judging witnesses by their manner, or counting how often a name recurs in the evidence. Over time, accurate assessments and arguments accumulate on the table alongside mistakes and misreadings. Yet despite this multitude of contributing voices, the relay jury is, in the end, forced to deliver a single verdict with one voice.

Does the jury *understand* the case? After all, not a single juror had a comprehensive grasp of the entire case. The verdict was not the direct expression of a single locus of understanding. It emerged from a complex interplay between hundreds of voices of uneven quality. The question this raises is not the familiar one from social epistemology – what to say when many minds coalesce into one collective agent (List and Pettit 2011; Bird 2010) – but rather its mirror image:

what to say when what *presents* as a unified individual mind disaggregates, upon inspection, into many voices or sub-personal mechanisms.

The relay jury offers a useful analogy for thinking about understanding in large language models (LLMs). LLMs process information through successions of layers that each deploy their own collection of parallel mechanisms. These mechanisms read from the communal table – the model’s *residual stream*. Each mechanism looks for different things and writes its own conclusions back into the stream for later layers to read. Interpretability research has shown these mechanisms to be uneven in quality and highly selective in what they attend to (see, e.g., Elhage et al. 2021; Lindsey et al. 2025). Like the jury room, the model must return a single verdict, which conceals the multitude of voices it contains. We call this condition *polyphony*, in an echo of the philosopher Mikhail Bakhtin’s account of Dostoevsky’s novels, which hold many independent voices and perspectives without subordinating them to a single authorial vision, leaving the whole to emerge from their interplay.¹

We contend that the polyphonic character of LLMs productively complicates the recent debate over whether such systems can be said to understand. The positions in that debate fall into three camps. One camp – call them the “anthropomorphisers” – unabashedly applies cognitive terms like “understanding” to LLMs.² A second camp – the “deflationists” – admonishes against the use of cognitive terms in this connection, maintaining that LLMs are best conceptualised in statistical or mathematical terms – we should speak of “matrix multiplications” producing outputs from “superficial correlations”, “distributional statistics”, or “pattern-matching”.³

Though it is more often practised than defended in print, there is a third position, which dismisses the entire debate as a distraction: as long as the models produce the outputs we want, these “indifferentists” maintain, we should not worry too much about how they got there.⁴ Whether AI models “understand” is a philosophical question in the pejorative sense of the term: an idle speculation.

We argue that the question of AI understanding is anything but idle – it *needs* to be raised even by the most hard-nosed practitioners, because it *guides the allocation of trust*. Yet making headway with it requires breaking out of what we call the *monophonic paradigm*: the perennially tempting idea that a cognitive system’s understanding of a subject matter must be localised to one substrate, a *locus of understanding*, which simultaneously underpins all the capacities that this understanding confers and reliably finds expression in them.

¹ See also Shaw (2026), whose formally oriented account of “polyphonic intelligence” provides independent evidence of the fruitfulness of the metaphor. Whereas Shaw develops a theoretical proposal for non-dominating integration among plural inferential processes, we draw on mechanistic findings about existing LLMs to explore the implications of polyphony for attributions of understanding.

² See (Piantadosi and Hill 2022; Bubeck et al. 2023; Søgaard 2023; Mandelkern and Linzen 2024; Brooks et al. 2024; Cappelen and Dever 2025).

³ See Chomsky, Roberts, and Watumull 2023, Weber 2024, Yiu, Kosoy, & Gopnik 2023, Bender, Gebru, McMillan-Major, & Shmitchell 2021, Marcus (2018), Floridi (2023), Bishop (2021), Yiu, Kosoy, & Gopnik (2023), and Bender and Hanna (2025).

⁴ See (Krishnan 2020; Carlini & Scarfe 2025).

This monophonic paradigm works well enough in its home territory of human cognition. Thinking of many capacities as tracing back to one seat of understanding helpfully licenses the largely reliable inference from an individual's display of competence in one respect – in *explaining* something, for instance – to the presumption that this individual will prove competent in other respects as well, i.e. in controlling, diagnosing, predicting, or reasoning counterfactually about that thing. Conversely, failure at any of these is reasonably taken to count against the presence of understanding more broadly (e.g. “if you can't explain it, you don't really understand it”).

When transposed to LLMs, however, the monophonic paradigm leads us astray, because the inferences it underwrites no longer reliably hold. Conversing with an LLM may *feel* like interacting with a unified mind that speaks with a single voice, but mechanistic inspection reveals that LLMs are in fact massively polyphonic systems. What we would ordinarily treat as a single concept may be represented in several partly overlapping ways; conversely, the same part of the model may contribute to several different representations or tasks. At the level of procedure, multiple pathways may have to combine to produce an answer, while the same task may also be discharged by alternative, sometimes overlapping combinations of pathways. Their contributions can reinforce, inhibit, replace, or compensate for one another. Every word an LLM produces reverberates with many voices.

We argue that a conception of understanding capable of making sense of LLMs must be a correspondingly polyphonic one. We prefer “polyphonic” to “distributed” (Hutchins 1995), because distributed cognition involves spreading a task outward across many agents and artifacts, whereas polyphony involves a spreading inwards, a subdivision into multiple mechanisms that look, from the outside, like a single agent. Nor is the relevant form of understanding best described as “fragmented” or “fractured”, as in Freeborn's (2026) account of AI understanding. This evokes a lost unity – yet polyphony is not necessarily a defect. The relay-jury structure is its own kind of unity. It constitutes a distinct form of understanding, and can also be a source of robustness and accuracy.

A significant consequence of polyphony, however, is that an outwardly competent performance may be underpinned by a multitude of criss-crossing mechanisms, none of which rise to a level we would ordinarily want to dignify with the term “understanding”. Conversely, a disappointing performance may conceal truly ingenious and sound procedures whose contributions get drowned out by competing signals. The monophonic paradigm, which encourages us to treat competent behaviour as testimony to a single underlying locus of understanding, licenses a number of inferences that become hazardous in the face of such polyphonic systems. As Millièrre and Rathkopf (2026) argue, this is a form of *anthropocentric bias*: we transfer to AI models a conception of understanding that is keyed to human cognition, and are led astray by insufficient appreciation of how profoundly these systems can differ from us.

It will be objected that human cognition is itself not monophonic: the brain does many things at once, and while philosophy has tended to emphasise the *unity* of consciousness (Bayne 2010), theories of cognition have found uses for pictures of a divided mind, from Minsky's (1986) society of mind through Dennett's (1991) multiple drafts to Kahneman's (2011) fast and slow

systems.⁵ But the point is that the monophonic paradigm works in practice for human cognition. Whether through internal coordination, external scaffolding, or a mixture of both, we have learned to fashion human individuals into suitable units for attributions of understanding – beings in whom the relevant capacities keep company reliably enough to treat them as ramifications of a single seat of understanding. This fashioning of a polyphonic interior into a stable unit has not (yet) happened for LLMs to the same extent. This is one reason why the inherited concept of understanding does not quite fit them.⁶

Here, we propose to remedy this lack of fit by re-engineering the concept of understanding for LLMs in light of the mechanistic evidence of their polyphony. We thereby distance ourselves from all three camps in the AI understanding debate. Against the indifferentists, we argue that the question whether AI understands is anything but idle: it does need to be asked, and a great deal turns on the answer; while it is true that one's answer depends on one's conception of understanding, the conclusion to draw is that the debate needs to ascend to the metaconceptual level and face the question of *how best to conceptualise* understanding in the context of AI systems. Against the deflationists, we maintain that the cognitive register, and some conception of understanding in particular, is *indispensable*: as our *reintroduction argument* establishes, mathematical and statistical vocabulary is too indiscriminate to separate trustworthy from untrustworthy outputs, and any vocabulary rich enough to do so ends up reintroducing the concept of understanding in all but name. Against the anthropomorphisers, we note that our existing concept of understanding is *inadequate*, because it carries monophonic assumptions that do not transfer to LLMs. We thus need to *adapt* the way we conceptualise understanding. Drawing on a range of recent findings in mechanistic interpretability research, we show how a suitably adjusted conception of understanding can gain a foothold in the actual workings of LLMs and do much-needed work in our interactions with AI.

We proceed as follows: §2 shows why we *need* a coherent way to think about understanding in LLMs if we are to know when to trust their outputs. §3 lays out what it would mean for mechanistic organisation to vindicate such trust. As §4 then brings out, however, LLM understanding is pervasively polyphonic, which complicates attributions of understanding in several respects. §5 then addresses how these complications play out at the scale of frontier models and synthesises the main ways in which we can determine whether to trust AI outputs. §6 concludes by considering the prospects for orchestrating polyphony.

2. The Reintroduction Argument

The question of how to conceptualise AI cognition is not merely a verbal dispute. In a growing range of situations, people must decide whether and how far to trust AI outputs. We contend that they cannot do so without forming a view of what the model has and has not understood.

⁵ See Evans & Stanovich (2013) for an overview of dual-process theories of higher cognition.

⁶ Though we focus on understanding, polyphony also complicates other cognitive attributions at the level of the model as a whole, including attributions of belief, intention, and knowledge.

Consider a doctor consulting an AI model about a perplexing case. Suppose the model proposes an intriguing diagnosis. The doctor must then decide how far to trust it. For this, she needs to know whether the model has *understood* the disease: whether it has grasped how the disease's causes, symptoms, and treatments hang together, or whether it is merely parroting medical language. The concept of understanding does indispensable work here: it guides the allocation of epistemic trust.

The doctor's predicament is far from unique. A judge weighing AI-generated sentencing recommendations must discriminate between recommendations grounded in a grasp of how precedents bear on the present case and ones that just pattern-match based on surface features; a professor deciding whether to direct students to an AI tutor must determine whether the model's command of the material warrants an endorsement. These practitioners cannot afford to shirk these discriminative tasks by dismissing the understanding question as idle philosophising. They have a pressing *conceptual need* – an instrumental need for a conceptual resource capable of doing the work traditionally performed by the concept of understanding. And if that work essentially involves differentiating *between* AI systems and their outputs, this also means that the question cannot remain whether the inherited concept of understanding applies to AI systems *as a class*. The question must be how our inherited concept of understanding needs to be adapted to perform discriminative labour *within* that class.

This is why indifferentism is unsustainable as a response to the understanding question. The question is anything but idle: it *demand*s an answer because a lot turns on it. And the point is not that *theorists* need to address it; it is *practitioners* – doctors, judges, and teachers – that need to address it. The question is ineluctable *in practice*.

Now deflationism, on the other hand, requires a more complex response, because deflationists may grant the ineluctability of the question, but insist that with LLMs, it uniformly receives a negative answer: LLMs can *never* properly be said to “understand.” Indeed, one can exhaustively describe what they do without employing any cognitive vocabulary. A transformer converts its input into numerical “embeddings”, passes them through layers of linear projections and non-linear transformations, and returns a probability distribution over possible next tokens. An LLM is thus a complicated function from input text to a probability distribution over continuations.

Such mathematical descriptions are perfectly correct, as far as they go. The problem is that they do not go far enough: they are too indiscriminating. They apply equally whether the model gives an accurate diagnosis or a ludicrous one – whether it solves a problem using a robust and general procedure or using a fragile heuristic. Mathematical descriptions apply across the cases we need to distinguish, without supplying the nuances required to tell those cases apart. Even a complete mathematical specification of a forward pass through the model would not, by itself, tell us whether the model is matching superficial linguistic patterns or deploying anything like joint-carving concepts organised to mirror the structure of a domain. Yet that is the question that matters. And any serious attempt to address is bound to reintroduce the concept of understanding in all but name.

Consider the doctor again. Even if she takes the deflationist's advice and tries to think about her predicament entirely in non-cognitive terms – purely in terms of pattern-matching, for instance – she must decide whether to trust the model's output. And this forces a discrimination between different *kinds* of pattern-matching. Not all pattern-matching is equal: some forms merit deference; others do not. But what *is* the kind of pattern-matching that merits deference? It is, one might think, pattern-matching that *tracks the underlying structure of the domain* – notably, by being sensitive to the causal dependencies between aetiological factors, pathogenetic mechanisms, and clinical manifestations rather than just to superficial correlations between them.

The deflationist can of course speak of “structure-tracking pattern-matching” and link that to epistemic trustworthiness; but the *distinction* she thereby draws, and the inferential implication she ties it up with, already begins to reproduce the inferential profile of the understanding/parroting distinction. The concept of understanding is being reintroduced through the back door.

The heart of what we call *the reintroduction argument*, then, is that any attempt to draw the distinction between reliable and unreliable AI outputs in deflationary vocabulary will end up reforging just the inferential connections that are characteristic of the vocabulary one seeks to avoid, thereby effectively reintroducing the concept of understanding, or at least a recognisable descendent of it, in all but name.

The fundamental reason for this is that non-cognitive vocabulary is *expressively inadequate* to the task of discriminating between the finer shades of competence and reliability that we need to distinguish in dealing with systems as sophisticated as contemporary AI models.

To see this, consider what happens if our doctor recognises the need for finer conceptual discrimination, but remains resolved to avoid cognitive vocabulary. She could look to *track records of success*. These are readily available as benchmark statistics: a 95% score on a medical-diagnosis benchmark (such as CPC-Bench, MSDiagnosis, or the Sequential Diagnosis Benchmark) would seem to offer some grounds for trust.

Yet success on benchmarks remains compatible with extensive reliance on shortcuts and superficial cues. A track record of success is too coarse an indicator to distinguish between a trustworthy and an untrustworthy model with 95% score. Even if that track record extends across a diversity of benchmarks and datasets, a gap remains between good performance on tests and real-world reliability on a novel case.

Our doctor needs to know whether the model can be trusted on *this particular real-world case*. She may therefore want to consider the model's real-world track record across cases that are relevantly similar to the case at hand. Yet even supposing real-world data to be available, this begs the question of which cases to *count* as relevantly similar. It was because the case was unusual and unobvious that she consulted the model to begin with; and *which* cases are relevantly similar to this one depends on the causal structure of the disease at hand – which is precisely what she is trying to find out in asking the model, and thus cannot help her in deciding whether to trust the model's answer.

What these attempts to evaluate a model using its performance history are fundamentally missing is sensitivity to the *internal basis* of the model's performance. Only something that *persists* across cases can underwrite trust in the model's competence on a novel case. One thus needs to look at what goes on inside the model – at the output-producing *process*.

Yet the moment one tries to identify trustworthiness-conferring properties of that process, one starts to recapitulate the literature's attempts to articulate what understanding is. Suppose the doctor asks whether the model is sensitive to the domain's *difference-makers* – whether its outputs covary with what an intervention on the disease would change. That is the mark of understanding on causal accounts, which characterise understanding as a grasp of what-would-happen-if relationships and of the factors on which an outcome depends (Pearl 2018; Woodward 2003; Strevens 2008). Suppose she asks whether the model's success is *robust* – surviving structure-preserving variations and extending to relevantly novel cases. That is understanding as cognitive control over a subject matter: the capacity to apply, adjust, and redeploy what one grasps beyond the cases in which one learned it (Hills 2016; de Regt 2017). Suppose she asks whether the model's outputs on related questions *hang together* – whether its diagnosis coheres with what it says about the disease's course, complications, and treatment. That is understanding as grasping connections (Wittgenstein 1953; Riggs 2003; Grimm 2011; Elgin 2017; Kvanvig 2018). Or suppose she asks whether the model has *compressed* its cases into an underlying rule. That is understanding as unification (Friedman 1974; Kitcher 1989; Schurz and Lambert 1994; Wilkenfeld 2018; Freeborn 2026).

The list could be extended, but the reintroduction argument concerns the pattern it instantiates. The literature on understanding maps out the inferential terrain characteristic of the concept. And while accounts differ over which region of the terrain is central, the point is that any discriminative resource that rises to the doctor's challenge must occupy some region of that terrain; it may do so without using the word "understanding", but it will at least partly reproduce the inferential role occupied by the word. And whether or not we take inferential role to *individuate* concepts, these inferential connections are what the debate over AI understanding is about; a deflationism that reintroduces those connections has effectively conceded the debate.

The cognitive register is thus hardly an indulgence; it is the register in which the relevant facts come into view.

3. Circuitry that Licenses Trust

What anthropomorphisers neglect, however, is that a concept can be indispensable and still unfit as it stands. The inherited concept earns its extension to AI through the important inferential connections it draws. But it also carries inferential connections that reflect its human origins. Re-engineering the concept for AI is a matter of sorting out which connections need *preserving* because they continue to do indispensable work, which need *forging* so that they anchor the concept in the distinctive make-up of models, and which need *severing* because they do not transfer to AI. Millièrè and Rathkopf anticipate that investigating LLM capacities will yield

“a novel ontology of cognitive kinds, optimized for explaining the distinctive strengths and weaknesses of machine intelligence rather than human intelligence” (2026, 386). What follows is an attempt to supply one entry in that novel ontology.

If attributions of understanding to AI are to be more than behaviourist courtesy, they must gain a foothold in what models actually contain. Following Beckmann and Quelo (2026), we take the mechanistic basis for LLM understanding to comprise three kinds of structure: *features*, *connections* between features, and *circuits* encoding more or less general procedures.

LLM layers read from and write to a shared internal state – the residual stream. That state is a long list of numbers that can be pictured as a latent *space*, in which each number configuration picks out a point. As the numbers change, that point shifts in a corresponding direction.

It appears that LLMs encode concepts by letting each direction in the space stand for something: they treat the directions as representing the degree to which something exhibits a certain feature. Researchers identified a direction that stands for the presence of the Golden Gate Bridge across languages, paraphrases, and imagery, for example (Templeton et al. 2024). These internal representations have become known as “features.” To demonstrate their causal efficacy, features can be *steered*: if one moves the point picked out by the list of numbers further along the direction representing the bridge, the LLM starts fixating on the bridge.

To learn facts about the world, an LLM needs to learn to *connect* its features in the right way. Connections between features take the form of learned dependencies by which activating one feature drives up (or down) the activation of another. This enables a model to represent the fact that there is a connection between Michael Jordan and basketball, for example, which allows it to retrieve the fact that Michael Jordan is a basketball player (Figure 1).

A *circuit*, finally, is a set of model components, typically scattered across several layers, that are wired together to implement some less content-specific *procedure* instead of relying on memorised facts. A well-understood example is a circuit formed to compute modular addition (Nanda et al. 2023a; Chughtai et al. 2023; Li et al. 2025; Beckmann and Quelo (2026)).




Figure 1. Connections between features ensure that the model completes the sequence “Michael Jordan plays” with “basketball”. The main role of attention heads is to carry forward features between residual streams (a). MLPs mainly combine features and recall features via activated features (b). Features are activated via directions in the residual stream’s latent space.

We propose to use the term *sound circuitry* to refer to a constellation of features, connections, and circuits that licenses trust in an AI model – relative to a certain type of task in a certain domain and across a certain range of cases. Such a constellation is *sound* to a degree insofar as, under an independently warranted interpretation of its causally efficacious states, those states bear a non-accidental, approximately structure-preserving relation to the elements and dependencies of the domain that matter to the task, and its transitions exploit that relation so that, were the constellation recruited and allowed to govern the computation, it would meet the correctness standard with corresponding reliability across the specified scope. *Soundness* is thus a graded term of art rather than the logician’s all-or-nothing notion.

The conception of AI understanding we want to develop centres on sound circuitry as what gives the notion of understanding a grip on the internal organisation of LLMs. Let us work through two case studies to substantiate this claim.

3.1 Connecting Features for Medical Diagnosis

Lindsey et al. (2025) examined what goes on under the hood when an LLM performs a differential diagnosis. They gave Claude 3.5 Haiku the prompt: “A 32-year-old female at 30 weeks gestation presents with severe right upper quadrant pain, mild headache, and nausea. BP is 162/98 mmHg, and labs show mildly elevated liver enzymes. If we can only ask about one other symptom, we should ask whether she’s experiencing...” – and Claude provided “visual disturbances” as the likeliest completion, followed by “proteinuria” (too much protein in urine).

To a doctor, these are the right continuations. The question is how Claude gets there. Does it rely on co-occurrence patterns between words in the prompt and those continuations? Or does it do what a doctor does, i.e. recognise the symptoms as typical of *preeclampsia* and then reason that hitherto unmentioned symptoms of preeclampsia include visual disturbances and proteinuria? The direct route requires no understanding of what is being diagnosed; the indirect route requires a grasp of the connections between the symptoms presented, the disease, and its further symptoms.

To determine which route Claude uses, one needs to know not only which features are active in the model, but what the connections between these features are, and which features activate which. There is a technique for exposing this internal wiring. It involves producing a so-called *attribution graph* (Ameisen et al. 2025; Lindsey et al. 2025).

One difficulty in producing such a graph is that almost every neuron fires for a miscellany of things – it is *polysemantic*. This is because the network needs to represent far more concepts than it has neurons by packing them in *superposition* (Olah et al. 2020; Elhage et al. 2022). To nonetheless home in on the features that are key to generating an output, one trains a second network – a *transcoder* – to reproduce the first network’s computation using replacement units engineered so that only a small fraction are active at any moment and each tends to be interpretable as encoding one concept. This gives us a sparse representation of the features that are active in this one inference.

But this does not yet tell us how these features *hang together*, i.e. which features activate which. Because the transcoder’s features enter the computation additively, one can calculate how much each feature *directly contributes* to activating later features. By representing each feature as a node and each contribution as a line or “edge”, one can then generate an attribution graph showing which features activated which on the path from prompt to output.

Initially, this graph depicts millions of edges. It needs pruning to a subgraph that preserves most of the computation. This gives us a sparse representation of the *connections between features* involved in the computation.⁷

An attribution graph can reveal whether the model took the superficial or the understanding-like route. If the model merely parroted medical text based on superficial statistics, the edges should run straight from the features for the prompt’s words to the feature for “visual disturbances”. On the understanding-like route, the edges should pass through diagnostic features such as the *preeclampsia* feature. It turns out that Claude indeed takes that latter route (Figure 2).

⁷ The graph reconstructs the model only approximately, however. What it cannot capture gets bundled into uninterpretable “error” terms.




Figure 2. Attribution graph for Claude 3.5 Haiku (Lindsey et al. 2025). Patient-status features activate a *preeclampsia* feature alongside competing hypotheses (*biliary system disorders*); the *preeclampsia* feature in turn activates features for confirmatory symptoms, which inform the model’s leading completions: “visual disturbances” and the runner-up, “proteinuria”.

To test whether the *preeclampsia* feature is causally efficacious in producing the output rather than merely correlated with it, the experimenters inhibited that feature. The model’s leading recommendation then flipped to asking about “decreased appetite” – a sign of biliary system disorders, the model’s alternative diagnosis. Not only has the model learned to wire up the connections between features to mirror the actual connections between preeclampsia and its symptoms; the model also brings features for rival hypotheses into play, much as a differential diagnosis would.

This is merely a look under the hood of a single inference, of course. Its ability to ground epistemic trust is correspondingly limited. It can warrant *local* trust, for this type of diagnosis in that type of case, by showing that the answer issues from the right kind of connections between features. But these connections are content-specific. To warrant broader trust, a different kind of internal structure is needed.

3.2 General Addition Circuits

What could warrant broader trust is a more general, more content-independent circuit. One of the first clear examples of this turned up in a small neural network trained to perform modular addition, i.e. addition with a ceiling (the *modulus*) after which one starts over. Trained on a large table of examples, the network seemed to just memorise those. But when a researcher went on vacation and left the model training, they found on their return that accuracy on unseen examples had unexpectedly shot up after all. The model had also become internally simpler, replacing its memorised examples with a compact computational circuit (Power et al. 2022).

Reverse-engineering this circuit, Nanda et al. (2023a) found that the model represents each number as an angle on a circle and adds these angles, so that the reset at the modulus happens automatically when coming full circle. This neat trick requires converting the landing point back into a number, however, which involves locating the peak of a relatively flat curve. To find this peak reliably, the model performs addition on multiple circles rotating at different speeds. Faster rotation sharpens the peak, yet also creates several of them. The model resolves that ambiguity by combining the curves – this suppresses false peaks while amplifying the one on which all curves agree.

One might suspect such clever tricks to emerge only when training a small network on nothing but modular addition. In fact, however, the trick of representing numbers on several circles at once was found in larger, generalist models as well – and even when computing *regular* addition (Nikankin et al. 2025; Kantamneni and Tegmark 2025).

Llama-3.1-8B offers a particularly striking example of a general circuit along these lines. It was shown to reuse the same compact circuitry across addition-like tasks (Figure 3) – not just regular addition, but also “What month is six months after August?” or “What day is five days after Wednesday?” (Feucht et al. 2026). This is surprising, since models are known to represent cyclic concepts such as months, weekdays, and 24-hour times using dedicated circular representations with 12, 7, or 24 positions (Zhou et al. 2024; Kantamneni and Tegmark 2025; Feucht et al. 2026). Researchers expected the model to exploit these dedicated representations also when *adding* months, weekdays, or 24-hour times. Computing “four months after October” on a circle with a period of 12 – that is, twelve positions – would, after all, make the *modulo* 12 operation automatic.

Yet this is not what Feucht et al. (2026) found. Llama-3.1-8B performs all kinds of additions – whether involving plain arithmetic or months, weekdays, or hours – in the same place: a circuit of just 28 neurons inside the MLP at layer 18. Cyclic concepts like months get translated into ordinary numbers and sent through a base-10 “calculator”. Thus, “What month is eight months after June?” becomes $8 + 6$. The calculator uses circles for this, but not those with periods of 12, 7, and 24 that would be natural for months, weekdays, and hours. Instead, it uses circles with periods of 2, 5, 10, 20, 50, and 100. This decimal machinery can handle any kind of addition. But it, too, uses multiple circles in parallel to resolve ambiguities: on the period-5 circle, 14, the sum of $8 + 6$, is indistinguishable from 4, 9, and 19; it takes the period-20 circle to pin it down. In later layers, a *modulo* 12 operation then maps the result – 14 – back to months, yielding the correct answer: February.




Figure 3. Instead of using separate calculators for each cyclic domain such as months, weekdays, or 24-hour times, Llama-3.1-8B reuses a single base-10 addition mechanism across addition-like tasks (Feucht et al. 2026). This mechanism consists of several circles with periods of 2, 5, 10, 20, 50, and 100.

The discovery of such a circuit implementing a truly general and reliable algorithm supports trusting the model beyond this one case: insofar as that circuit continues to be recruited and to control outputs, Llama-3.1-8B can be expected to handle unseen examples across the various guises that addition takes. This is illustrative of the kind of broader trust that a general and reliable circuit can underwrite.

These examples of sound circuitry help substantiate the idea that attributions of understanding to AI models are, at bottom, claims about internal organisation. To say that a model *understands* is to claim that, *were* we to reverse-engineer what features, connections, and circuits are causally responsible for its performance, our findings would vindicate trust. The diagnosis and addition cases show that mechanistic interpretability can in principle uncover such organisation.

Yet the convenient singulars in which we described these findings – *the* preeclampsia feature, *the* addition circuit – conceal something important. Talk of “*the*” *preeclampsia* feature is itself shorthand, used by researchers to group several internal representations that encode subtly different contents, but play roughly the same role in the computation – a complication registered by the stacked boxes in Fig. 2. The general addition circuit likewise subdivides into multiple parallel computations that are ambiguous in isolation, but unambiguously point to an answer when considered together. In both cases, the singular label suggests a single, self-contained mechanism where we actually find an interacting plurality. This prefigures the argument of the next section: that LLMs are pervasively polyphonic.

4. Polyphony

Polyphony is the condition in which a model’s performance is enabled by a motley mix of mechanisms of uneven reliability, whose contributions may combine, compete, or compensate for one another. In §1, we introduced the notion through the jury room analogy. We now

develop that analogy into a working model of transformer computation (§4.1) and assemble the mechanistic evidence that LLMs are polyphonic at the level of features, connections between features, and circuits (§4.2). We then show how this organisation disrupts the inference patterns surrounding the concept of understanding (§4.3), before gathering the resulting conditions into a polyphonic conception of AI understanding (§4.4).

4.1 The Jury Room as a Model of Transformer Computation

The relay jury introduced in §1 can now be mapped more precisely onto a transformer. Each morning brings a fresh panel – one layer of the network – which inherits only the accumulated notes on the communal table. The table is the residual stream; its notes are the features currently active there.

Each juror first selects material according to a fixed rule or perspective, attending, for example, to causal relations, numerical regularities, contextual cues, or keywords. This is the role of attention heads: they identify relevant information from a particular perspective, while ignoring much of the rest and often overlapping in what they select.

The jurors then work in parallel, combining their selections, retrieving background information, flagging inconsistencies, or merely copying salient details forward. Some contributions are sophisticated; others are superficial, heuristic, or mistaken. All are added to the table for the next panel. This processing stage corresponds to the multilayer perceptrons (MLPs), which transform and enrich the selected information.

Across successive days, lines of thought accumulate, reinforce one another, and sometimes interfere. A group of jurors collaborating across days to execute a procedure corresponds to a circuit, though any juror may contribute to several circuits, just as an attention head or neuron may do.

Significantly, no juror sees the complete picture. What understanding of the case emerges resides in the cumulative interplay of all these parallel contributions. No presiding juror was designated at the beginning, so where there is coordination among the contributors, it arises *sua sponte*.

Yet by the final day, the jury must deliver a verdict with one voice. That verdict corresponds to the next token prediction. It may be that this verdict coincides with the one a lone expert with a comprehensive grasp of the case would have reached. But the relay jury's best work cannot be isolated from the surrounding tangle of simpler inferences and heuristics, because they are held up and enabled and shaped by them. The verdict is irreducibly a collective achievement.

Table 1 summarises the correspondences.

Jury room element	LLM component
Communal table	Residual stream
A note on the table	A feature
Each day and panel of jurors	One layer

Jury room element	LLM component
New jury each day	No memory between layers (except what is in the residual stream)
Selecting evidence	Attention heads
Processing evidence, recalling information	MLPs
A group of jurors performing one procedure across several days	A circuit
Final verdict	Model output (next predicted token)

Table 1. Summary of the jury room analogy.

4.2 The Mechanistic Case for Polyphony

The jury room analogy gives intuitive form to what is, at bottom, a mechanistic thesis: LLMs are polyphonic at several levels. A single content may be carried by multiple features; within a circuit, multiple pathways may make complementary, redundant, or opposing contributions; and distinct circuits may perform the same task (Fig. 4).




Figure 4: Polyphony at three levels: (a) *features*, where related features jointly encode a concept; (b) *connections*, where direct, indirect, and shortcut routes among features converge on the same output; and (c) *circuits*, where two low-overlap combinations of features and connections are each sufficient to implement the procedure, and one expands its contribution when the other is ablated.

First, there is polyphony at the level of *features*. What a human would subsume under a single concept often appears in LLMs as a whole family of cognate features playing broadly the same computational role. As with the *preeclampsia* feature, what a tidied-up attribution graph collects under the heading of the *Texas* feature actually presents as multiple *Texas*-related features that *collectively* perform the role of activating “the” *Texas* feature (Lindsey et al. 2025). Researchers also find features recurring at successive stages of processing and in states associated with neighbouring prompt tokens (Lindsey et al. 2024; Lindsey et al. 2025). Such recurrence may reflect real duplication, or an important signal being carried forward, or a

concept being differentiated into more fine-grained ones. Talk of “the” feature for X is thus often shorthand for multiple representations that either redundantly successively, or jointly encode one concept. Since this accumulation of features risks filling up the residual stream, some attention heads and MLPs act as cleanup mechanisms, cancelling selected activations by writing their negations into the stream (Elhage et al. 2021).

Second, there is polyphony at the level of connections, or *intra-circuit polyphony*. A circuit may comprise several pathways, each consisting of a different sequence of connections between features and carrying out part of the overall computation. When Claude 3.5 Haiku completes “The capital of the state containing Dallas is” with “Austin”, for example, Dallas-related features activate Texas-related features, which combine with features directing the model to name a capital. But there is also a shortcut from Dallas to Austin; and both the Texas and “say a capital” clusters affect the output directly as well as indirectly via a further “say Austin” cluster (Lindsey et al. 2025). One feature can thus influence the output through several convergent routes.

Another example of intra-circuit polyphony is the way Claude performs two-digit addition using a coalition of heuristics (Ameisen et al. 2025; Lindsey et al. 2025). For $36 + 59$, one pathway uses stored associations between addition problems and their answers to estimate that the sum is somewhere around 92, while another considers only the final digits to infer that the answer ends in 5. Only once these intersect (“a number near 92, ending in 5”) does the correct answer (95) emerge. Each supplies a constraint that determines the answer only when conjoined with the other. Accordingly, these are best considered not distinct circuits, but complementary pathways within a *coalitional computation*.

Intra-circuit polyphony enables pathways to compensate for one another – a phenomenon that McGrath et al. (2023) have dubbed “the hydra effect”. To complete “When John and Mary went to the shop, John gave a drink to ...”, GPT-2 Small must retrieve “Mary”. Attention heads copy the appropriate name into the output prediction. If one head is suppressed, backup heads in later layers compensate (Wang et al. 2023; McDougall et al. 2024).

Third, there is also abundant evidence of *inter-circuit polyphony* – individually sufficient circuits performing the same task. In their aptly titled “All Circuits Lead to Rome” paper, Chen et al. (2026) offer the clearest challenge to the idea that LLM understanding must have one privileged internal home. They search GPT-2 for a small subgraph that can sustain performance on a task when all other connections are masked, then repeat the search while penalising the reuse of connections. Were there such a thing as *the* circuit underpinning the model’s performance, the searches should converge on it. Yet across twenty searches, they find multiple circuits with little overlap. Even within an apparent three-connection core, no connection is indispensable: exclude any one, and the wider model supplies another sufficient route. Each discovered circuit is therefore only one of a larger repertoire of sufficient mechanisms. This undermines the monophonic expectation that successful performance must ultimately rest on a single privileged mechanism.

Polyphony at the level of circuits also creates possibilities for *destructive interference* between circuits. Kim et al. (2025) and Valentino et al. (2025) identified a circuit that correctly evaluates

sylogistic inferences based on logical form, but whose voice gets drowned out by the contributions of content-based mechanisms. The model is less likely to approve the syllogism: “All apples are vegetations. All vegetations are institutions. Some apples are institutions” than: “All apples are edible fruits. All edible fruits are fruits. All apples are fruits”. Both are valid. But because the former contains substantively implausible premises, the content-based circuits interfere with the formal circuit’s correct verdict.

Because there is polyphony both within and between circuits, a successful computation can be *compositionally polyphonic*, *realisationally polyphonic*, or both. It is compositionally polyphonic when no one contribution is sufficient and several must operate as a coalition. It is *realisationally polyphonic* when it is realised by more than one circuit, so suppressing one leaves others able to take over. And it is both at when it performs a task using multiple alternative coalitions that are each internally dependent on a plurality of pathways.

There is, then, plenty of mechanistic evidence of polyphony in transformer-based LLMs. Indeed, the transformer architecture may itself be conducive to polyphony. Attention heads and MLPs write additively to a shared residual stream, so their contributions remain available to downstream components unless actively negated (Elhage et al. 2021). Multiple partial, redundant, or conflicting signals can therefore coexist and interact without being reconciled by a central coordinating mechanism. Though not imposed by the architecture, polyphony is given room to develop in it. And one would expect the wide variety of tasks that generalist models are trained to favour its emergence.

4.3 How Polyphony Complicates Attributions of Understanding

The pervasiveness of polyphony means that we cannot count on identifying mechanisms that are consistently involved whenever a model performs a certain task. We may wish to identify *the* mechanism, since that would give us an intellectual and practical handle on these forbiddingly complex systems. Yet even when their internal organisation becomes interpretable, it resists reduction to a single mechanism.

For the debate over AI understanding, the significance of polyphony lies in how it disrupts inferences that our inherited concept of understanding encourages us to draw – inferences that are defeasible already in the human context, but that become downright hazardous in pervasively polyphonic systems like LLMs.

Consider the inference from the failure to perform to the absence of understanding. There are several reasons why sound circuitry may be present without finding outward expression.

A first possibility is *dormancy*: suitable circuitry is present, but not activated. Nikankin et al. (2025) find that Llama-3 relies on a collection of arithmetic heuristics that fire only over a restricted range of operands. Dormancy can also be purposely induced. Adding a suffix to a dangerous prompt can inhibit the activation of a feature that would ordinarily recruit a refusal mechanism (Arditi et al. 2024; Ball, Kreuter, and Panickssery 2026).

A second possibility is what might be called *misrecruitment*: sound circuitry is active, but the mechanisms that situate a novel case within it are too closely tailored to the training regime. “TaxiGPT” illustrates this. Vafa et al. (2024) trained a GPT-2-style transformer from scratch to predict a Manhattan taxi’s next turn from its origin, destination, and preceding turns (“N NW NE E SW...”). Although it produced legal turns nearly all the time, it occasionally implied physically impossible street configurations – such as streets labelled NW but facing east – or jumps over intervening streets. Its performance also deteriorated sharply when repeated forced detours led away from the destination. Vafa et al. concluded that the model was “very far from recovering the true street map of New York City” (2024, 2).

Yet Beckmann, Quelo, and Freitas’s (2026) mechanistic diagnosis of TaxiGPT’s failures undercuts this inference from behavioural error to a lack of spatial understanding. They recovered a fully accurate map of Manhattan from TaxiGPT’s activations and traced the model’s errors to a problem arising downstream of recovering a map. On in-distribution tests, the model was able to correctly locate itself on the map. But on harder tests – when the goal was twice as distant as during training, or when the taxi was forced by researchers to turn away from its destination two thirds of the time – the activation written to the correct location can be too weak. As noise accumulates across the map features, a wrong intersection can become active and produce an illegal move. This is *misrecruitment* rather than dormancy: an accurate map is present and recruited, but the mechanism locating the taxi within it does not generalise robustly out of distribution. The problem is thus not that TaxiGPT lacks a correct map of Manhattan. The map circuitry lies ready to be recruited, but because the localisation mechanism fails on out-of-distribution cases, it is not always recruited correctly.

A third possible problem is *interference*: relevant circuitry is activated, but its contribution is dominated by less reliable mechanisms. This is what happened with the circuit evaluating the formal validity of syllogisms. Rai et al. (2025) similarly find that language models can err on parentheses-balancing tasks when unreliable mechanisms swamp more reliable ones. Amplifying the reliable mechanisms can restore accuracy from approximately zero to nearly 100%. The original failure therefore did not show that the model lacked sound circuitry. It was present and active, but outvoted.

Millière and Rathkopf (2026) identify a fourth possible problem: *compute bottlenecks* at inference-time. A transformer required to answer in a single forward pass is limited in what it can compute (Merrill and Sabharwal 2024). A failure may therefore reflect insufficient computational room to use and coordinate the model’s existing resources. Increasing a model’s budget of thinking tokens can lift this constraint and remedy the failure (Wei et al. 2023; Muennighoff et al. 2025).

In the terms of the jury-room analogy, the expert juror may be missing altogether (no sound circuitry), or present but silent (dormancy), bring their expertise to bear on a garbled characterisation of the case (misrecruitment), speak but be outvoted (interference), or lack enough time to complete and coordinate the relevant reasoning (compute bottleneck). A

mistaken verdict cannot tell us which of these conditions obtained and therefore does not establish that the relevant understanding was absent.

Yet the converse inference – from successful performance to a general and reliable basis for it – is rendered equally hazardous by polyphony. Eshuijs, Wang, and Fokkens (2025) showed this by fine-tuning GPT-2 on movie reviews in which particular actor names were correlated with positive or negative sentiment. The model classified reviews reliably when the review sentiment matched the association attached to the actor mentioned. Yet mechanistic analysis revealed that attention heads used the actor’s name to steer the verdict before even fully processing the review. Substituting a name with the opposite association often flipped the classification. Successful performance may thus rest on a shortcut whose unreliability emerges only once the incidental correlation is broken.

Polyphony also weakens the correlation and integration between the different capacities conferred by understanding. With humans, there is an imperfect, but noteworthy correlation between the capacities to *perform*, *explain*, *predict*, *diagnose errors in*, and *reason counterfactually* about a task one understands. A person’s ability to explain the task or diagnose mistakes is therefore reasonably treated as evidence that they can apply the relevant principles themselves. Moreover, the principles they invoke in explaining the task normally also guide their own performance. The relation among these capacities is thus not merely an accidental correlation: the capacities are intelligibly coordinated. Call this expected correlation and coordination among the capacities associated with understanding *cross-capacity integration*.

These capacities should be expected to come apart more readily in pervasively polyphonic systems, since the different capacities may be supported by completely distinct and isolated circuitry. Claude’s polyphonic addition strategy offers an example: asked how it arrived at 95 when adding $36 + 59$, Claude invokes the schoolbook procedure of adding the ones, carrying, and then adding the tens. That is a sound method for performing the calculation, but mechanistic analysis shows that this was not how Claude performed it. The capacities to perform and explain the calculation are both displayed, but they are not integrated.

Mancoridis et al. (2025) provide a still more striking dissociation. Asked what an ABAB rhyme scheme is, GPT-4o correctly explains that the first and third lines must rhyme, as must the second and fourth. Yet when prompted to complete an ABAB poem whose first line ends with “out”, it ends the third line with “soft”. Asked whether “out” rhymes with “soft”, the model correctly answers that it does not. The relevant principles are therefore available for explanation and error diagnosis without acquiring sufficient control over generation. The display of one capacity associated with understanding thus provides no guarantee that the others are present.⁸

Mancoridis et al. interpret this to mean that the LLM offers only an *illusion* of understanding – *Potemkin understanding*, as they call it, in reference to the fake village fronts that Russian

⁸ Dehghanighobadi et al. (2025) provide another example.

minister Grigory Potemkin is said to have erected along the Dnieper river to impress Catherine the Great. But it may be more accurate to interpret this as an example of *polyphonic* understanding. After all, there is not *nothing* behind the facade: the model was able to explain the rhyme scheme and diagnose errors in its implementation. The behavioural evidence available cannot tell us whether the appropriate circuitry is completely absent, or present but unrecruited, or recruited but misapplied, or correctly applied but overridden.

4.4 A Polyphonic Conception of AI Understanding

Polyphony does not make attributions of understanding impossible, but it complicates them. A conception of understanding fitted to polyphonic systems must be sensitive to these complications. Recall what the concept of understanding needs to do for us: it must help practitioners distinguish outputs that merit trust from those that do not. This requires a *modal* conception, attuned not merely to what the model has done in the past, but to the internal organisation that persists within it, making it *likely* to perform well on future cases.

Section 3 distinguished three kinds of internal structure on which such sound circuitry may draw: features that track elements of a domain, connections that encode relevant dependencies among them, and circuits that implement procedures for exploiting that structure. In a polyphonic system, these components need not constitute a single, self-contained mechanism. A task may be performed by a coalition whose members are individually insufficient, and the model may contain several alternative coalitions capable of performing it. What matters is whether some such constellation bears a non-accidental, approximately structure-preserving relation to the aspects of the domain relevant to the task and would, if correctly recruited and governing the computation, meet an independently specified standard of correctness across the stated range of cases.

In light of the complications presented in 4.3, we can derive four desiderata on such sound circuitry. The first and most obvious one is that it must be *present*. The model must contain some combination of features, connections, and circuits that preserves, to the degree required by the task, the domain distinctions and dependencies on which correct performance counterfactually depends. This requirement does not privilege one realisation. A model may contain several combinations that answer adequately to the domain, overlapping to a greater or lesser extent. Presence requires that at least one sound combination be available, not that a particular component be shared by every successful performance.

Second, sound circuitry must be *recruited*: it must be activated by the case at hand rather than lying dormant. A model may contain accurate structure that is reliably engaged across familiar cases but not by inputs outside the range on which it was trained.

Third, as the TaxiGPT example showed, sound circuitry must be recruited *correctly*. If the mechanisms mediating its application to a novel case are insufficiently reliable, this vitiates the output even if sound circuitry is present and active.

Fourth, the correctly recruited circuitry must be *in control* of the output. Its causal contribution must carry through to what the model says or does rather than being overwhelmed by competing mechanisms. Control need not take the monophonic form of one mechanism prevailing alone. The governing circuitry may be a coalition whose contributions are jointly sufficient under the relevant background conditions even though no contribution is sufficient by itself. Another coalition may take over if one pathway is suppressed. Membership in the coalition that produces an output on one occasion is therefore not a mark of indispensability: a different combination might have delivered the same output by another route. What the control condition excludes is not causal plurality, but cases in which the contribution of sound circuitry is outweighed by unsound competitors.

These nuances allow us to separate the *depth* of a model's understanding from its *hold* over behaviour. Depth concerns the richness, accuracy, and generality of the sound structure present in the model: how much of the domain has it grasped, and at what level? Hold concerns whether that structure is reliably recruited and in control: how consistently is what the model has grasped reflected in its outputs? A model may contain deep understanding that rarely controls its answers, or shallow understanding that controls them consistently.

Behavioural evidence struggles to distinguish depth from hold. Shallow understanding may produce the same correct output as deep understanding; indeed, a model may articulate a deeply principled procedure without ever relying on it itself. This is how mere fluency comes to be overread as understanding. Conversely, deep understanding without a strong hold over outputs may produce the same fragility as the absence of deep understanding, which is how fragility comes to be overread as evidence of shallowness.

The preceding discussion yields a conception of AI understanding re-engineered for polyphonic systems. On this conception, understanding is always of a domain, with respect to a task or capacity, to some degree, and across some scope. Once these dimensions are fixed, the definition can be stated succinctly:

AI understanding. A model understands if (i) it contains circuitry that is sound to the specified degree for the relevant domain and task or capacity; (ii) it reliably and correctly recruits some such circuitry across the specified scope; and (iii) the recruited circuitry reliably exercises sufficient control over the relevant outputs within that scope.

Evidence that these conditions are met provides defeasible, *pro tanto* warrant for relying on those outputs. This conception thus preserves the connection that §2 showed to be indispensable: attributions of understanding guide the allocation of epistemic trust. It forges links to features, connections, and circuits that give the concept a foothold in the internal organisation of LLMs.

What this conception relinquishes is the monophonic presumption that understanding must have one privileged internal home, be manifested uniformly across the capacities associated with it, or be transparently reflected in behaviour. Nothing in the definition requires a unique internal realiser. The relevant circuitry may be multiply realised within the model. What matters

is not that it always recruits the same mechanism, but that it possesses a persistent capacity to recruit some sound circuitry and give it sufficient hold over its behaviour.

Polyphony therefore does not close off the question of understanding. It multiplies the questions that must be answered while making them tractable in principle through mechanistic investigation.

5. Assessing Understanding in Frontier Models

The evidence we considered came from relatively small models, and one might object that this tells us little about the frontier models. Those systems are vastly larger, which suggests two reasons for pessimism: they have more space to memorise what they are trained on, and they are too complex to reverse-engineer in full.

While the difficulty is real, it does not render our proposed conception of AI understanding idle. Evidence for the presence, recruitment, and control of sound circuitry admits of degrees, and it can be assembled even here.

The findings from smaller models have a limited, but important role in this. For a given task and domain, they establish that transformer architectures *can* form features tracking domain structure, connect them in appropriate ways, and combine them into procedures that generalise beyond memorised examples. These findings act as possibility proofs. If smaller models can do it, larger models ought to be able to do it, too.

In fact, evidence suggests that larger models tend, if anything, to generalise even better than smaller ones. The “lottery ticket hypothesis” surmised that since a bigger network contains more sub-networks, it is likelier to harbour a winner (Frankle & Carbin 2019). Yet newer research offers more systematic explanations. One is that larger models are less likely to get trapped in a local minimum, because the added width also widens traps in the loss landscape, opening escape routes towards better solutions (Simsek et al. 2021; Martinelli, Brea and Gerstner 2026). Another is that larger models are more likely to land on simpler or more compressible solutions, because these take up a larger share of the space (Wilson 2025).⁹ A third explanation is that the extra capacity lets models form circuitry for rare tasks, which are harder to learn because they provide weaker training signals, and which a smaller model has to neglect in favour of frequent tasks (Huang et al. 2026).

These considerations weaken the first reason for pessimism: scale need not favour memorisation over generalisation. But being better equipped to form general procedures is not the same as actually forming and appropriately recruiting, and relying on them. The second reason for pessimism is therefore more pressing.

⁹ Among the weight-settings that fit the training data, some sit on narrow spikes, and others in broad basins, where a whole neighbourhood of settings does equally well. The basins constitute more compressible solutions, since there is no need to specify them precisely; and since volume compounds across dimensions, every parameter added increases the compressible solutions’ share of the space.

Yet full reconstruction is not necessary for every attribution of understanding. A targeted investigation can establish how a model performs a specified task over a bounded range. What is required is enough evidence to assess the conditions identified in §4:

1. *Soundness and presence.* Can the relevant domain structure be decoded from the model's internal states, and is that structure causally exploited rather than merely correlated with the output? Do the dependencies among the features preserve the task-relevant dependencies in the domain? Can researchers identify a circuit, or a family of alternative and overlapping coalitions, whose operation meets an independently specified standard of correctness?
2. *Recruitment.* Is some sound constellation activated across the cases for which understanding is attributed, or does suitable circuitry remain dormant outside a narrow range of inputs or prompts? Is it correctly recruited even in challengingly novel cases?
3. *Control.* When sound circuitry is recruited, does its contribution govern the output, or is it outweighed by competing mechanisms? Interventions can test whether amplifying, suppressing, or redirecting a pathway changes the result as predicted. This can also expose compensating mechanisms.
4. *Scope and integration.* Across which inputs, contexts, distribution shifts, and capacities do presence, recruitment, and control continue to obtain? Does the model's competence extend across application, explanation, prediction, error diagnosis, and counterfactual reasoning?

Where mechanistic access remains limited, behavioural evidence can nevertheless bear on the same conditions, though less conclusively. Relevant indicators include generalisation to appropriately novel cases, robustness under immaterial changes in framing, sensitivity to changes that are structurally relevant, and controlled out-of-distribution tests that map where performance breaks down. Tests across clusters of capacities likewise offer clues concerning cross-capacity integration.

But the system cards of leading labs are increasingly publishing mechanistic evidence to show that relevant features were identified and shown to predictably alter behaviour. This points to a future where system cards might publicly certify and advertise a model's governance by sound circuitry for certain tasks and domains. We are used to testing human understanding through exams and collections of hard problems. Yet the study of human understanding has not had the benefit of the tools and possibilities that interpretability research now has at its disposal. With the repertoire of white-box methods growing at pace, we should expect the study of AI understanding to push far beyond behavioural evaluations, in ways that human-facing attributions of understanding never could.

6. The Orchestration Gap

The question whether an LLM understands is unavoidable but, in its familiar form, poorly posed. It is unavoidable because practitioners must decide which outputs merit epistemic trust, and purely mathematical or statistical descriptions cannot by themselves draw the distinctions that

decision requires. The reintroduction argument showed that any vocabulary rich enough to distinguish trustworthy from untrustworthy outputs reinstates at least part of the inferential role of understanding. Yet the familiar monophonic paradigm directs us to seek a single locus of understanding that supports application, explanation, prediction, diagnosis, and counterfactual reasoning alike, and manifests itself transparently in behaviour. This may be a tolerable idealisation of human understanding, but it becomes a liability for systems as polyphonic as current LLMs. In such systems, success need not disclose a reliable basis, failure need not show that no such basis exists, and competence in one capacity need not extend to others. Behavioural fluency and fragility do not interpret themselves.

The positive proposal of this paper is a conception of AI understanding centred on sound circuitry. A model understands something when it contains some constellation of features, connections, and circuits that bears a non-accidental, approximately structure-preserving relation to the task-relevant organisation of the domain, and that would meet an independently specified standard if correctly recruited and allowed to govern the computation. Any such attribution must therefore specify a task and scope; within that scope, the relevant circuitry must be present, reliably and correctly recruited, and in control. A model's depth of understanding concerns the richness, accuracy, and generality of its sound circuitry; its hold concerns how reliably that circuitry is recruited and governs what the model says or does. Because sound circuitry may consist of a coalition of individually insufficient components, and because several alternative coalitions may perform the same task, understanding need have no privileged internal home.

Polyphony is not peculiar to LLMs. Neuroscience calls the capacity of structurally different elements to perform the same function "degeneracy" (Edelman & Gally, 2001). In the brain, degeneracy allows distinct neural pathways to sustain the same cognitive capacity and confers robustness when one route fails (Price & Friston, 2002; Noppeney et al., 2004). Its familiar signature is *graceful degradation*, which LLMs also exhibit: sizeable blocks of layers can be ablated with surprisingly modest loss (Michel et al., 2019; Gromov et al., 2024). Yet the epistemological literature still gravitates towards pictures of understanding as something paradigmatically unified and integrated.

The two pictures can be reconciled if unity is treated not as the standing architecture of understanding, but as an achievement of orchestration. People confident that they understand a familiar mechanism quickly discover how fragmentary their understanding really is when asked to explain it (Rozenblit and Keil 2002). Yet, given room to work things out, they can assemble these fragments into the unified form that philosophical accounts describe. What remains of Kepler's laws years after a class may be only a set of half-connected fragments, resembling what Freeborn (2026) calls "fractured understanding" in deep nets. Asked to use those laws to help one's children predict a planet's position, however, one may reconstruct the connections with pen and paper, testing and revising until the pieces cohere. Nor is this peculiar to the layperson: Boge, Schuster, and Stoll (2026) argue that even a scientist's grasp of a phenomenon is often only a tacit model, and counts as *scientific* understanding only once it coalesces into an explanation that the scientific community can follow. On this picture, unified

understanding emerges when partial resources are successfully coordinated, often with social and material scaffolding.

There is emerging evidence that scaffolding can similarly improve orchestration in LLMs. External scratchpads and agentic harnesses provide a workspace in which different parts of a model can be successively recruited, tested, and recombined. Venhoff et al. (2026) offer a striking example. They found that the chain-of-thought of a fully trained reasoning model decomposes into a small repertoire of distinct moves, such as planning, checking, and backtracking. Interestingly, those activation signatures were already present in the base model. And a simple trained classifier could select which move should come next and nudge the base model's activations accordingly; this intervention closed most of the gap between the two models. What the base model lacked was not the capacity to make the individual moves, but reliable control over when to deploy them. This is the *orchestration gap*: a failure to coordinate competences that the system already possesses.

Current LLMs nevertheless do not orchestrate their partial competences as reliably as humans.¹⁰ Until this gap closes, practitioners cannot rely on the familiar inferences associated with the monophonic conception of understanding. But they can rely on those licensed by a polyphonic conception, which turns the elusive question of AI understanding into tractable empirical questions about the presence, recruitment, and control of sound circuitry – questions open to mechanistic investigation and intervention. If the relay jury's verdict conceals the polyphony that produced it, we must look to the notes on the table to tell whether that verdict deserves our trust.

¹⁰ As illustrated by the alien failure modes LLMs continue to display (e.g., McCoy et al., 2024; Chen et al., 2025).